

Topics: Permutation tests, logistic regression

Due on November 16, 2019

Name of student: SUBHRAJYOTY ROY

Roll number: MB1911

1. Permutation test of independence

[8]

Perform permutation tests of independence on Table 1 using the χ^2 and the likelihood ratio test statistics. Plot the permutation null distributions of these statistics. Check how the results depend on the number of permutations used. Compare the results with the standard χ^2 and likelihood ratio tests.

Gender	Handedness			Total by Gender
	Right-handed	Left-handed	Ambidextrous	
Male	14	3	3	20
Female	13	1	2	16
Other	6	1	1	8
Total by Handedness	33	5	6	44

Table 1: Gender and Handedness

Solution. While χ^2 and Likelihood ratio test statistic can be computed by hand for the given table, I use the following R code to read the data into R and perform the test to circumvent unnecessary calculation.

```

1 data <- matrix(c(14, 3, 3, 13, 1, 2, 6, 1, 1), nrow = 3, byrow = T, dimnames =
  list("Gender" = c("Male", "Female", "Other"), "Handedness" = c("Right", "Left", "
  Both")))
2 print(data)
3 chisq.test(data)

```

Gender	Handedness		
	Right	Left	Both
Male	14	3	3
Female	13	1	2
Other	6	1	1

Pearson's Chi-squared test

```
data: data
X-squared = 0.81, df = 4, p-value = 0.9371
```

Since the p-value is 0.9371, which is extremely higher than the significance level of 0.05, it is clear that this data does not show any evidence that Gender and Handedness are not independent, thereby enabling us believe in favour of independence between Gender and Handedness.

```
1 DescTools::GTest(data)
```

Log likelihood ratio (G-test) test of independence without correction

```
data: data
G = 0.85984, X-squared df = 4, p-value = 0.9303
```

The above code snippet shows the output of Likelihood ratio test of independence. Clearly, the p-value obtained is quite similar to that of the usual χ^2 test statistic, hence showing no evidence in rejecting in favour of the belief that Gender and Handedness are not independent.

To perform the permutation test, the dataset is first reshaped into a desired view.

```

1 meltdata <- reshape2::melt(data)
2 meltdata <- as.data.frame(lapply(meltdata, rep, meltdata$value))[, -3]
3 meltdata$Gender <- factor(meltdata$Gender)
4 meltdata$Handedness <- factor(meltdata$Handedness)
5 meltdata

```

	Gender	Handedness
1	Male	Right
2	Male	Right
3	Male	Right
...		
14	Male	Right
15	Female	Right
16	Female	Right
...		
27	Female	Right
28	Other	Right
...		
33	Other	Right
34	Male	Left
35	Male	Left
36	Male	Left
37	Female	Left
38	Other	Left
39	Male	Both
40	Male	Both
41	Male	Both
42	Female	Both
43	Female	Both
44	Other	Both

Since there are lots of permutations ($\frac{44!}{20!16!8!} \approx 1.29 \times 10^{18}$) of the **Gender** variable, therefore, it may not be possible to completely enumerate all possible permutations of the **Gender** variables keeping the other variable fixed.

We begin by writing two customized function in **R**, by which we can compute the Pearson's chi-square test statistic and Likelihood Ratio test statistic according to the following equations.

$$\chi^2 = \sum_i \sum_j \frac{(n_{ij} - E_{ij})^2}{E_{ij}}$$

and

$$-2 \log \Lambda = 2 \sum_i \sum_j n_{ij} \log \left(\frac{n_{ij}}{E_{ij}} \right) \text{ where } E_{ij} = \frac{n_{i.} n_{.j}}{n}$$

```

1 chisq.stat <- function(tab) {
2   n <- sum(tab)
3   rowmars <- apply(tab, 1, sum)
4   colmars <- apply(tab, 2, sum)
5   exp.tab <- rowmars %*% t(colmars)/n
6
7   return(sum((tab - exp.tab)^2 / exp.tab))
8 }
9
10 lrt.stat <- function(tab) {
11   n <- sum(tab)
12   rowmars <- apply(tab, 1, sum)
13   colmars <- apply(tab, 2, sum)
14   exp.tab <- rowmars %*% t(colmars)/n
15

```

```

16   return(2*sum(tab * log(tab / exp.tab), na.rm = T)) # if one cell becomes 0,
17     log is not defined, so we omit them in the sum
18 }

```

We then proceed to write a customized function in **R**, where we consider some permutations of the **Gender** variable. Under the null hypothesis that **Gender** and **Handedness** are independent of each other, it is clear that replacing the permuted row as the **Gender** variable, keeping the **Handedness** variable fixed, would give rise to a contingency table generated by the same mechanism as the original one. From that table, we can obtain the value of Pearson's chi square test statistic or an Likelihood Ratio statistic using above functions, and use that to infer about the p-value of the observed table.

```

1  perm.stat <- function(data, stat, nperms = 10e4, seed = 1234) {
2
3    if (! (stat %in% c("chisq", "lrt"))) {
4      stop("Not Implemented yet! Only 'chisq' and 'lrt' is available.")
5    }
6    else {
7      set.seed(seed)
8      x <- sapply((1:nperms), function(a){ sample(data[, 1], size = nrow(data)) })
9      x <- t(x)
10     x <- unique(x)
11
12     if (stat == "chisq") {
13       y <- apply(x, 1, function(a) {
14         tab <- table(a, data[, 2])
15         return(chisq.stat(tab))
16       })
17     }
18     else if (stat == "lrt") {
19       y <- apply(x, 1, function(a) {
20         tab <- table(a, data[, 2])
21         return(lrt.stat(tab))
22       })
23     }
24
25     return(y)
26   }
27 }

```

Now, we perform the Permutation test with chi-square test statistic with 5000 permutations, and obtain a p-value of 0.9298.

```

1  y <- perm.stat(meltdata, stat = "chisq", nperms = 5000, seed = 1911)
2  head(y)
3  paste("Approximated p-Value is", sum(y > chisq.stat(data))/5000)

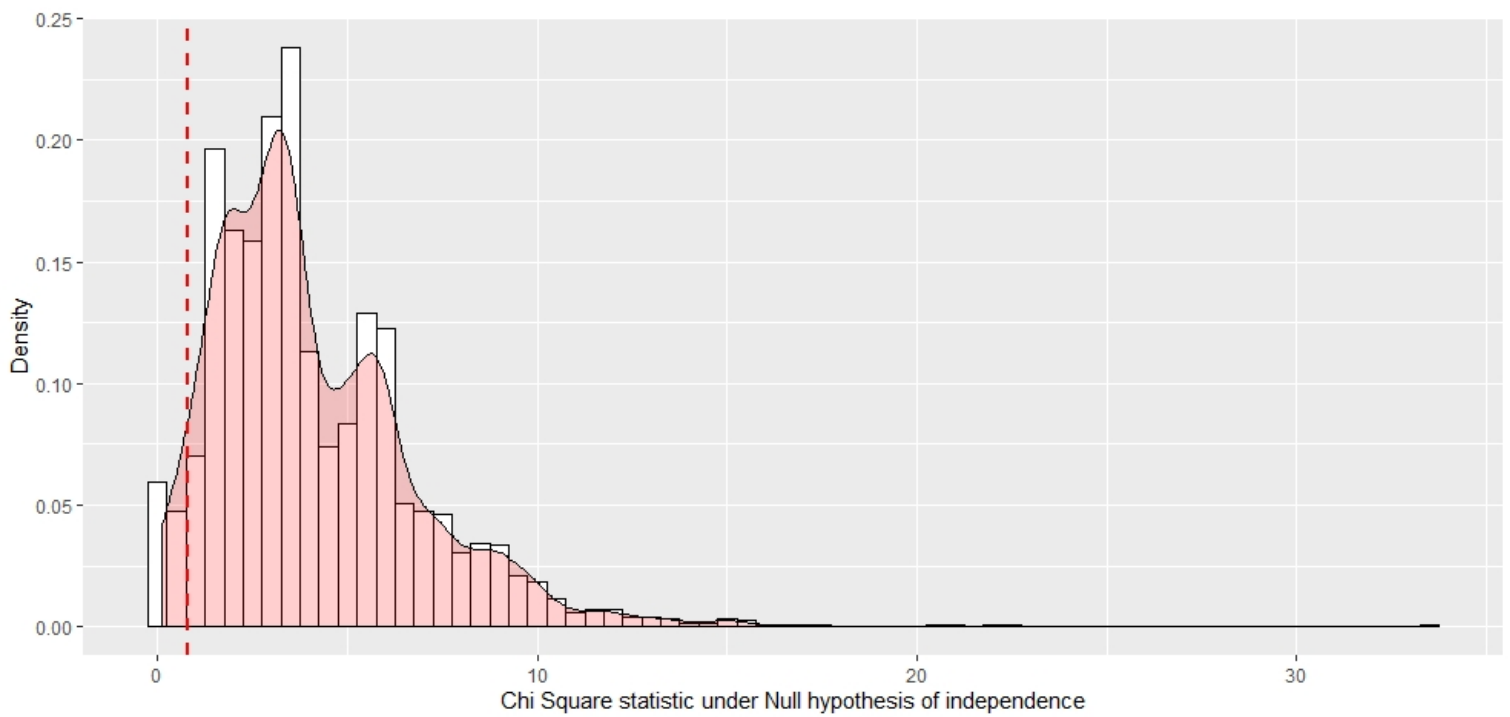
```

```

[1] 4.748333 5.348333 5.785000 2.451667 4.748333 2.785000
[1] "Approximated p-Value is 0.9298"

```

The plot for the null distribution looks as follows.

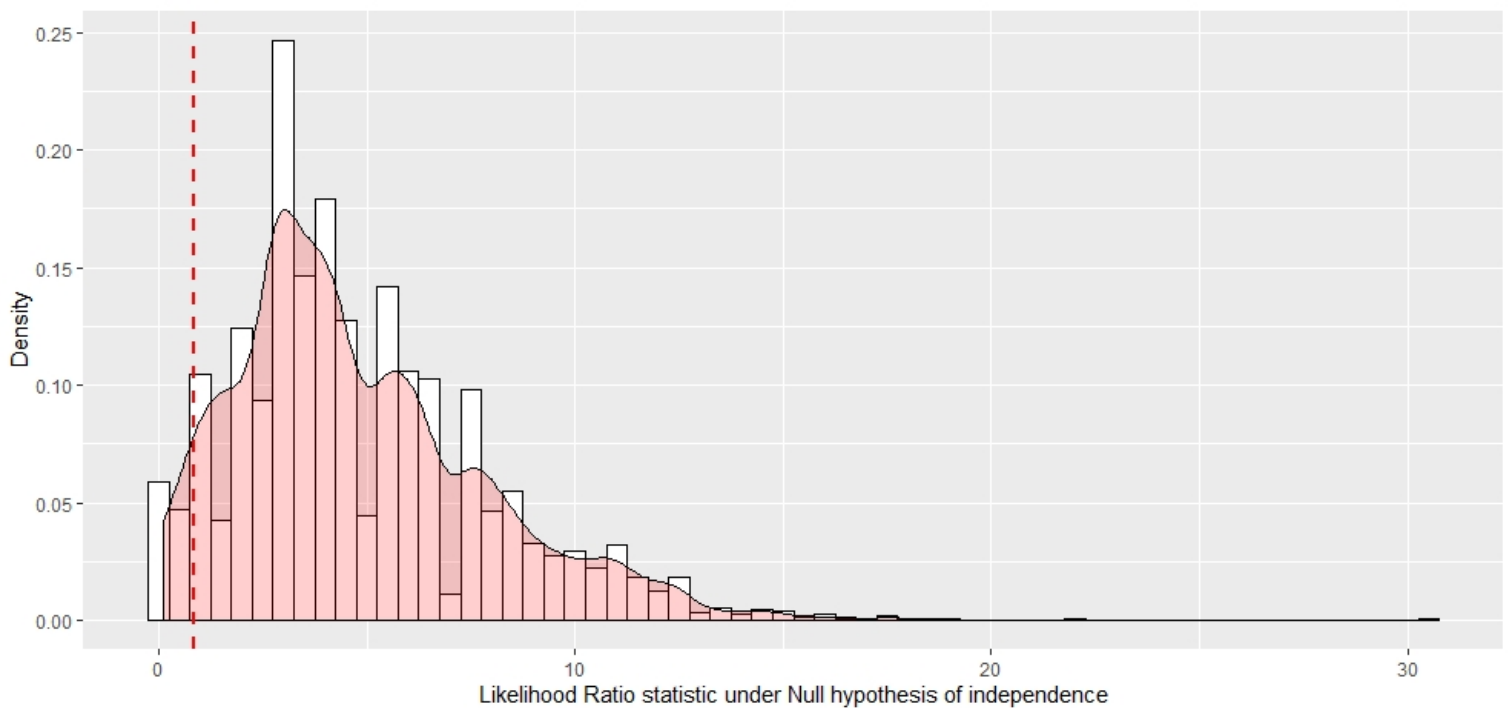


With same permutations, the p-value with Likelihood Ratio statistic turns to be exactly same, 0.9298, which should be true, as both the statistic measure the same kind of deviations from independence setup.

```
1 y <- perm.stat(meltdata, stat = "lrt", nperms = 5000, seed = 1911)
2 head(y)
3 paste("Approximated p-Value is", sum(y > lrt.stat(data))/5000)
```

```
[1] 5.655498 6.280661 5.001904 2.184919 4.975902 3.566309
[1] "Approximated p-Value is 0.9298"
```

and the null distribution looks as follows;



From the figures above, we find that the distribution of chi-square statistic is less dispersed than the null distribution of Likelihood ratio statistic. However, both the distributions are positively skewed, as expected and resembles the structure of a chi-square distribution.

Now, since both these p-values are 0.9298 which is extremely higher than the significance level of 0.05, hence we cannot reject the null hypothesis of independence between **Gender** and **Handedness**. As far as this data is concerned, it does not show any evidence that these two variables are associated in some way (i.e. not independent).

Also note that, the exact p-values obtained are extremely close to the approximate p-value based on asymptotic distribution of Pearson's chi-square and Likelihood Ratio test statistic. Clearly, it indicates that the asymptotic distribution well approximates the behavior of exact distribution in the light of the data.

Following is table containing the p-values for these permutation tests, based on the number of permutations used. All were produced with a seed of 1234 in the previous functions.

Number of Permutations	p-value with Chi-Square statistic	p-value with LRT
5	1	1
10	1	1
50	0.92	0.92
100	0.93	0.93
500	0.932	0.932
1000	0.932	0.932
5000	0.9324	0.9324
10,000	0.9336	0.9336
100,000	0.93652	0.93652

Interestingly, the p-values for both type of test statistic for permutation test turns out to be same.

2. Exact logistic regression

[12]

In a standard binomial logistic regression set-up

$$\text{logit}(\mathbb{P}(Y_i = 1)) = \beta_0 + \beta^\top X_i, \quad i = 1, \dots, n,$$

write down sufficient statistics T_j for the parameters β_j , $j = 0, \dots, p$. Show that the distribution of T_p conditional on T_0, \dots, T_{p-1} depends only on β_p . Thus, using this conditional distribution, one can estimate and perform inference on β_p . Does a conditional MLE always exist?

Write down the conditional distribution under $H_0 : \beta_p = 0$. Describe how you would do an exact test of this hypothesis.

Use the **R** packages **logistix** and **elrm** to perform exact logistic regression on the data in Table 2 with “White-collar job” as response, and “Gender” and “College education” as explanatory variables. Also, perform

Gender	College education	White-collar job	Number of cases
M	No	1	8
F	No	1	6
M	Yes	7	10
F	Yes	6	6

Table 2: Another fictitious dataset.

a standard logistic regression on the same data and compare the results.

Solution. We write the likelihood of the binomial logistic regression setup, where (X_i, Y_i) be the set of datapoints with Y_i taking values 0 and 1.

$$\begin{aligned}
L &= \prod_{i=1}^n (\mathbb{P}(Y_i = 1))^{Y_i} (\mathbb{P}(Y_i = 0))^{1-Y_i} \\
&= \prod_{i=1}^n \left(\frac{e^{\beta_0 + \beta^\top X_i}}{1 + e^{\beta_0 + \beta^\top X_i}} \right)^{Y_i} \left(\frac{1}{1 + e^{\beta_0 + \beta^\top X_i}} \right)^{1-Y_i} \\
&= \prod_{i=1}^n \left(\frac{1}{1 + e^{\beta_0 + \beta^\top X_i}} \right) \times e^{\beta_0 \sum_{i=1}^n Y_i + \sum_{j=1}^p \beta_j \sum_{i=1}^n X_{ij} Y_i}
\end{aligned}$$

where X_{ij} is the i -th observation corresponding to j -th variable (or predictor). Note that, the first term in the multiplication does not depend on the random observation Y_i , and hence is a fixed function of the parameters β , (as X_i 's are treated as constants). Therefore, by use of Neymann Factorization theorem, we obtain that; $T = (T_0, T_1, \dots, T_p)$ with $T_0 = \sum_{i=1}^n Y_i$ and $T_j = \sum_{i=1}^n X_{ij} Y_i$ is a sufficient statistic for $(\beta_0, \beta_1, \dots, \beta_p)$.

Note that, the joint density of (T_0, T_1, \dots, T_p) is; $f(t_0, t_1, \dots, t_p) \propto \exp \left[\sum_{j=0}^p \beta_j t_j \right]$. Therefore,

$$\begin{aligned}
f(t_p | t_0, t_1, \dots, t_{p-1}) &= \frac{f(t_0, t_1, \dots, t_p)}{f(t_0, t_1, \dots, t_{p-1})} \\
&= \frac{f(t_0, t_1, \dots, t_p)}{\sum_{t_p} f(t_0, t_1, \dots, t_p)} \\
&= \frac{c(t_0, \dots, t_p) \exp \left[\sum_{j=0}^p \beta_j t_j \right]}{\sum_{t_p} c(t_0, \dots, t_p) \exp \left[\sum_{j=0}^p \beta_j t_j \right]}, \text{ where } c(t_0, \dots, t_p) \text{ is the proportionality constant} \\
&= \frac{c(t_{(p-1)}, t_p) \exp [\beta_p t_p]}{\sum_{t_p} c(t_{(p-1)}, t_p) \exp [\beta_p t_p]} \text{ where } t_{(p-1)} = (t_0, t_1, \dots, t_{p-1})
\end{aligned}$$

which does only depend on β_p . Also note that, the conditional MLE does not always exist. This can be seen as follows. Since X_{ip} 's are fixed quantities, and Y_i is 0 or 1, it is clear that, the sum in the denominator runs through all elements of the set $\mathbb{T}_p = \{\sum_{i=1}^n Y_i X_{ip} : Y_i \in \{0, 1\}\}$. Let, the observed t_p^* that we get is $\max \mathbb{T}_p$, which exists since the set \mathbb{T}_p is finite for a fixed set of X_i 's. Then;

$$\begin{aligned}
f(t_p^* | t_0, t_1, \dots, t_{p-1}) &= \frac{c(t_{(p-1)}, t_p^*) \exp [\beta_p t_p^*]}{\sum_{t_p \in \mathbb{T}_p} c(t_{(p-1)}, t_p) \exp [\beta_p t_p]} \\
&= \frac{1}{1 + \sum_{t_p \in \mathbb{T}_p - \{t_p^*\}} \frac{c(t_{(p-1)}, t_p^*)}{c(t_{(p-1)}, t_p)} \exp [\beta_p (t_p - t_p^*)]} \\
&\rightarrow 1 \quad \text{as } \beta_p \rightarrow \infty, \text{ since } (t_p - t_p^*) < 0
\end{aligned}$$

Therefore, conditional MLE need not always exist in the above case. Also, note that the proportionality constants are positive. On the other hand, if the observed t_p^{**} that we obtain is actually $\min \mathbb{T}_p$, then;

$$\begin{aligned}
f(t_p^{**} | t_0, t_1, \dots, t_{p-1}) &= \frac{c(t_{(p-1)}, t_p^{**}) \exp [\beta_p t_p^{**}]}{\sum_{t_p \in \mathbb{T}_p} c(t_{(p-1)}, t_p) \exp [\beta_p t_p]} \\
&= \frac{1}{1 + \sum_{t_p \in \mathbb{T}_p - \{t_p^{**}\}} \frac{c(t_{(p-1)}, t_p^{**})}{c(t_{(p-1)}, t_p)} \exp [\beta_p (t_p - t_p^{**})]} \\
&\rightarrow 1 \quad \text{as } \beta_p \rightarrow -\infty, \text{ since } (t_p - t_p^{**}) > 0
\end{aligned}$$

Hence, in this case also, the conditional MLE need not exist.

Under the null hypothesis $H_0 : \beta_p = 0$, the conditional distribution reduces to; $f_{H_0}(t_p|t_0, \dots, t_{p-1}) = \frac{c(t_{(p-1)}, t_p)}{\sum_{u \in \mathbb{T}_p} c(t_{(p-1)}, u)}$.

Now, to perform a hypothesis testing for $\beta = 0$, note that; $f_{T_p}(t; \beta_p)/f_{T_p}(t; \beta_p = 0) \propto \exp[\beta_p t]$. As the likelihood ratio is monotonically increasing as a function of t , hence, we obtain the p-value for the observed value of the statistic $T_p = \sum_{i=1}^n Y_i X_{ip}$ as;

$$p_+ = \sum_{u \geq t} f_{T_p}(u; \beta_p = 0) = \frac{\sum_{u \geq t} c(t_{(p-1)}, u)}{\sum_{u \in \mathbb{T}_p} c(t_{(p-1)}, u)}$$

for testing against the alternative $\beta > 0$. Similarly, for testing against the alternative that $\beta < 0$, we could obtain the p-value as;

$$p_- = \sum_{u \leq t} f_{T_p}(u; \beta_p = 0) = \frac{\sum_{u \leq t} c(t_{(p-1)}, u)}{\sum_{u \in \mathbb{T}_p} c(t_{(p-1)}, u)}$$

And for testing against the two sided alternative; we can use the p-value as $2 \min\{p_+, p_-\}$. Then, we use this p-values and compare them with the nominal significance level α , to perform the hypothesis test (i.e. reject H_0 when the p-value is less than the significance level, and accept otherwise).

Next, we use package `elrm` to perform exact logistic regression on the given data.

```
1 df <- data.frame(Gender = c("M", "F", "M", "F"), CollegeEd = c("No", "No", "Yes", "Yes"),
2   WhiteJob = c(1, 1, 7, 6), Total = c(8, 6, 10, 6))
3
4 library(elrm)
5
6 # CI for Gender variable
7 fit <- elrm(WhiteJob / Total ~ Gender + CollegeEd, interest = ~ Gender, dataset
8   = df, iter = 10e5, burnIn = 100)
9 summary(fit)
10
11 # CI for College Education Variable
12 fit <- elrm(WhiteJob / Total ~ Gender + CollegeEd, interest = ~ CollegeEd,
   dataset = df,
   iter = 10e5, burnIn = 100)
summary(fit)
```

```
Progress: 100% Progress: 95% Progress: 90% Progress: 85% Progress: 80%
Generation of the Markov Chain required 35 secs
Conducting inference ...
Inference required 5 secs
```

Call:

```
elrm(formula = WhiteJob/Total ~ Gender + CollegeEd, interest = ~Gender, iter = 1e+06, dataset = df, burnIn
= 100)
```

Results:

	estimate	p-value	p-value_se	mc_size
GenderM	-1.35482	0.33699	0.00094	999900

95% Confidence Intervals for Parameters

	lower	upper
GenderM	-5.342778	1.128712

Call:

```
elrm(formula = WhiteJob/Total ~ Gender + CollegeEd, interest = ~CollegeEd, iter = 1e+06, dataset = df,
      burnIn = 100)
```

Results:

	estimate	p-value	p-value_se	mc_size
CollegeEdYes	3.6278	0.00031	2e-05	999900

95% Confidence Intervals for Parameters

	lower	upper
CollegeEdYes	1.107514	7.948637

From the above confidence interval, we note that the for the Gender variable, the 95% confidence interval does contain the value 0, and hence we are unable to reject that null hypothesis. So, conclude that effect of Gender on probability of getting White Collar Job is insignificant based on the data. However, for the College Education variable, the 95% exact confidence interval contains the value 0, and hence we can reject the null hypothesis that coefficient for College Education in logistic regression model is 0, and conclude that effect of College Education on probability of getting White Collar Job is significant at a 5% level of significance.

Using the package `logistiX`, almost similar confidence intervals are obtained;

```
1 library(logistiX)
2 Gender <- factor(c(rep("M", 8), rep("F", 6), rep("M", 10), rep("F", 6)))
3 CollegeEd <- factor(c(rep("No", 14), rep("Yes", 16)))
4 WhiteJob <- c(1, rep(0, 7), 1, rep(0, 5), rep(1, 7), rep(0, 3), rep(1, 6))
5 longdf <- data.frame(Gender = Gender, CollegeEd = CollegeEd, WhiteJob = WhiteJob
6                       )
7 x <- model.matrix(WhiteJob ~ Gender + CollegeEd, longdf)[, -1]
8 fit <- logistiX(x, y = longdf$WhiteJob)
9 summary(fit)
```

Exact logistic regression

Call:

```
logistiX(x = x, y = longdf$WhiteJob)
```

Estimation method:	LX
CI method:	exact
Test method:	TST

Summary of estimates, confidence intervals and parameter hypotheses tests:

	estimates	2.5 %	97.5 %	statistic	pvalue	cardinality
1	-1.360474	-5.367983	1.128986	8	0.4556514914	6
2	3.338699	1.101657	7.265954	13	0.0005817118	14

It seems that the `elrm` approximately obtains the same values of upper and lower bounds for confidence interval. Also, based on whether the confidence interval contains 0 or not, or checking whether the p-value is less than $\alpha = 0.05$ or not, we obtain same decision as before. The contribution of **Gender** variable is not significant in deciding the probability of getting a **White Collar Job**, but the contribution of **College Education** is significant.

We also perform the standard logistic regression on the data.

```
1 fit <- glm(cbind(WhiteJob, Total - WhiteJob) ~ Gender + CollegeEd, data = df,
2           family = "binomial")
3 summary(fit)
```



```
Call:
glm(formula = cbind(WhiteJob, Total - WhiteJob) ~ Gender + CollegeEd,
family = "binomial", data = df)
```

Deviance Residuals:

1	2	3	4
0.5628	-0.4437	-0.3179	0.9629

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.1471	0.8763	-1.309	0.19054
GenderM	-1.4515	1.2038	-1.206	0.22790
CollegeEdYes	3.6687	1.1904	3.082	0.00206 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 17.9365 on 3 degrees of freedom

Residual deviance: 1.5418 on 1 degrees of freedom

AIC: 13.877

Number of Fisher Scoring iterations: 4

We note that, using Wald's asymptotic test for standard logistic regression, we also find that Gender does not have a significant effect, whereas College Education has one on predicting the log odds of having White Collar job.